

ORIE 7391: Faster: Algorithmic Ideas for Speeding Up Optimization

Least squares

Professor Udell

Operations Research and Information Engineering
Cornell

January 31, 2022

solving a linear system

given **design matrix** $X \in \mathbf{R}^{m \times n}$, **righthand side** (rhs) $y \in \mathbf{R}^m$.
find w so that

$$Xw = y.$$

solving a linear system

given **design matrix** $X \in \mathbf{R}^{m \times n}$, **righthand side** (rhs) $y \in \mathbf{R}^m$.
find w so that

$$Xw = y.$$

how?

- ▶ factor and solve
 - ▶ QR
 - ▶ singular value decomposition (SVD)
 - ▶ Cholesky (for symmetric X)
- ▶ iterative methods
 - ▶ conjugate gradient (CG)
 - ▶ iterative refinement

we will talk about QR and CG

considerations in choosing a method

- ▶ one problem, or many righthand sides y with the same design matrix X ?
- ▶ sparse or dense X ?
- ▶ symmetric X or rectangular problem?

optimality condition for least squares is a linear system

given $X \in \mathbf{R}^{m \times n}$, $y \in \mathbf{R}^m$. find w to solve

$$\text{minimize} \quad \|Xw - y\|^2.$$

to solve, take gradient, set to 0. solution x satisfies **normal equations**

$$X^T X w = X^T y.$$

a linear system! (with symmetric positive semidefinite design matrix $X^T X$.)

Outline

QR

Conjugate gradient

Iterative refinement

The fundamental theorem of numerical analysis

Theorem

Never form the inverse (or pseudoinverse) of a matrix explicitly.

(Numerically unstable.)

Corollary: never type `inv(X'*X)` or `pinv(X'*X)` to solve the normal equations.

The fundamental theorem of numerical analysis

Theorem

Never form the inverse (or pseudoinverse) of a matrix explicitly.

(Numerically unstable.)

Corollary: never type `inv(X'*X)` or `pinv(X'*X)` to solve the normal equations.

Instead: compute the inverse using easier matrices to invert, like

- ▶ Orthogonal matrices Q :

$$a = Qb \iff Q^T a = b$$

- ▶ Triangular matrices R :

if $a = Rb$, can find b given R and a by solving sequence of simple, stable equations.

The QR factorization

every matrix X can be written using **QR decomposition** as $X = QR$

- ▶ $Q \in \mathbf{R}^{n \times d}$ has orthogonal columns: $Q^\top Q = I_d$
- ▶ $R \in \mathbf{R}^{d \times d}$ is upper triangular: $R_{ij} = 0$ for $i > j$
- ▶ diagonal of $R \in \mathbf{R}^{d \times d}$ is positive: $R_{ii} > 0$ for $i = 1, \dots, d$
- ▶ this factorization always exists and is unique
(proof by Gram-Schmidt construction)

can compute QR factorization of X in $2nd^2$ flops

The QR factorization

every matrix X can be written using **QR decomposition** as
 $X = QR$

- ▶ $Q \in \mathbf{R}^{n \times d}$ has orthogonal columns: $Q^\top Q = I_d$
- ▶ $R \in \mathbf{R}^{d \times d}$ is upper triangular: $R_{ij} = 0$ for $i > j$
- ▶ diagonal of $R \in \mathbf{R}^{d \times d}$ is positive: $R_{ii} > 0$ for $i = 1, \dots, d$
- ▶ this factorization always exists and is unique
(proof by Gram-Schmidt construction)

can compute QR factorization of X in $2nd^2$ flops

use `scipy.linalg.qr`:

$$Q, R = \mathbf{qr}(X)$$

advantage of QR: it's easy to invert R !

QR for least squares

use QR to solve least squares: if $X = QR$,

$$X^T X w = X^T y$$

$$(QR)^T QR w = (QR)^T y$$

$$R^T Q^T QR w = R^T Q^T y$$

$$R^T R w = R^T Q^T y$$

$$R w = Q^T y$$

$$w = R^{-1} Q^T y$$

Computational considerations

never form the inverse explicitly: numerically unstable!

instead, use QR factorization:

- ▶ compute QR factorization of X ($2nd^2$ flops)
- ▶ to compute $w = R^{-1}Q^\top y$
 - ▶ form $b = Q^\top y$ ($2nd$ flops)
 - ▶ compute $w = R^{-1}b$ by back-substitution (d^2 flops)

Computational considerations

never form the inverse explicitly: numerically unstable!

instead, use QR factorization:

- ▶ compute QR factorization of X ($2nd^2$ flops)
- ▶ to compute $w = R^{-1}Q^T y$
 - ▶ form $b = Q^T y$ ($2nd$ flops)
 - ▶ compute $w = R^{-1}b$ by back-substitution (d^2 flops)

in julia (or matlab), the **backslash operator** solves least-squares efficiently (usually, using QR)

$$w = X \setminus y$$

in python, use `numpy.lstsq`

Demo: QR

<https://github.com/ORIE4741/demos/QR.ipynb>

Sparse QR

complexity of QR depends on the sparsity of Q and R :

- ▶ compute QR factorization of X (?? flops)
- ▶ to compute $w = R^{-1}Q^T y$
 - ▶ form $b = Q^T y$ (**nnz**(Q) flops)
 - ▶ compute $w = R^{-1}b$ by back-substitution (**nnz**(R) flops)

Q-less QR

during QR, can compute $Q^T y$ essentially for free!

- ▶ compute QR of $[X \ y]$.

Q-less QR

during QR, can compute $Q^T y$ essentially for free!

- compute QR of $[X \ y]$.

or compute it afterwards without forming Q :

$$\begin{aligned} X^T b &= (QR)^T b = R^T Q^T b \\ R^{-1} X^T b &= Q^T b \end{aligned}$$

Cholesky and QR

consider **Gram matrix** $G = X^T X \succeq 0$. if $X = QR$,

$$G = R^T Q^T Q R = R^T R$$

this construction gives **Cholesky factorization**

- ▶ factors spd matrix into triangular matrices
- ▶ Cholesky factors of $X^T X$ have same structure as R

Sparse QR: exercise

- ▶ can you guess the sparsity of R given sparsity of X ?
- ▶ can you change sparsity of R by permuting columns of X ?

Sparse QR: exercise

- ▶ can you guess the sparsity of R given sparsity of X ?
- ▶ can you change sparsity of R by permuting columns of X ?

use 'colamd' in Matlab, equivalents in Python and julia

Chordal fill-in

to analyze fill-in

- ▶ consider spd matrix, for simplicity
- ▶ interpret matrix as directed graph
- ▶ form clique tree
- ▶ identify fill-in





Figure 4.1: *Left.* Filled graph with 9 vertices. The number next to each vertex is the index $\sigma^{-1}(v)$. *Right.* Array representation of the same graph.

Outline

QR

Conjugate gradient

Iterative refinement

Conjugate gradients

symmetric positive definite system of equations

$$Ax = b, \quad A \in \mathbf{R}^{n \times n}, \quad A = A^T \succeq 0$$

Conjugate gradients

symmetric positive definite system of equations

$$Ax = b, \quad A \in \mathbf{R}^{n \times n}, \quad A = A^T \succeq 0$$

why use conjugate gradients?

- ▶ uses only matrix-vector multiplies with A
 - ▶ useful for structured (from PDE or graph) or sparse matrices, easy to parallelize, ...
- ▶ most useful for problems with $n > 10^5$ or more
- ▶ converges exactly in n iterations
- ▶ converges approximately much faster
- ▶ quick-and-dirty solve is appropriate **inside** inner loop of optimization algo

Conjugate gradients

symmetric positive definite system of equations

$$Ax = b, \quad A \in \mathbf{R}^{n \times n}, \quad A = A^T \succeq 0$$

why use conjugate gradients?

- ▶ uses only matrix-vector multiplies with A
 - ▶ useful for structured (from PDE or graph) or sparse matrices, easy to parallelize, ...
- ▶ most useful for problems with $n > 10^5$ or more
- ▶ converges exactly in n iterations
- ▶ converges approximately much faster
- ▶ quick-and-dirty solve is appropriate **inside** inner loop of optimization algo

other variants for indefinite (MINRES) or nonsymmetric matrices (GMRES)

source: presentation of CG inspired by [https:](https://stanford.edu/class/ee364b/lectures/conj_grad_slides.pdf)

[//stanford.edu/class/ee364b/lectures/conj_grad_slides.pdf](https://stanford.edu/class/ee364b/lectures/conj_grad_slides.pdf)

Iterative methods for least squares

define

- ▶ objective $f(x) = \frac{1}{2} \|Ax - b\|^2$
- ▶ condition number $\kappa(A) = \sigma_n(A)/\sigma_1(A)$
- ▶ bound $R \geq \|x_\star\|$ on norm of solution x_\star
- ▶ goal: find apx solution within accuracy $f(x) - f(x_\star) \leq \epsilon$

Iterative methods for least squares

define

- ▶ objective $f(x) = \frac{1}{2} \|Ax - b\|^2$
- ▶ condition number $\kappa(A) = \sigma_n(A)/\sigma_1(A)$
- ▶ bound $R \geq \|x_\star\|$ on norm of solution x_\star
- ▶ goal: find apx solution within accuracy $f(x) - f(x_\star) \leq \epsilon$

how many iterations (matvecs) required?

Iterative methods for least squares

define

- ▶ objective $f(x) = \frac{1}{2} \|Ax - b\|^2$
- ▶ condition number $\kappa(A) = \sigma_n(A)/\sigma_1(A)$
- ▶ bound $R \geq \|x_\star\|$ on norm of solution x_\star
- ▶ goal: find apx solution within accuracy $f(x) - f(x_\star) \leq \epsilon$

how many iterations (matvecs) required?

- ▶ conjugate gradient
 - ▶ $O(\sqrt{\kappa} \log(\frac{1}{\epsilon}))$

Iterative methods for least squares

define

- ▶ objective $f(x) = \frac{1}{2} \|Ax - b\|^2$
- ▶ condition number $\kappa(A) = \sigma_n(A)/\sigma_1(A)$
- ▶ bound $R \geq \|x_\star\|$ on norm of solution x_\star
- ▶ goal: find apx solution within accuracy $f(x) - f(x_\star) \leq \epsilon$

how many iterations (matvecs) required?

- ▶ conjugate gradient
 - ▶ $O(\sqrt{\kappa} \log(\frac{1}{\epsilon}))$
- ▶ gradient descent (GD)
 - ▶ $O(\kappa \log(1/\epsilon))$

Iterative methods for least squares

define

- ▶ objective $f(x) = \frac{1}{2} \|Ax - b\|^2$
- ▶ condition number $\kappa(A) = \sigma_n(A)/\sigma_1(A)$
- ▶ bound $R \geq \|x_\star\|$ on norm of solution x_\star
- ▶ goal: find apx solution within accuracy $f(x) - f(x_\star) \leq \epsilon$

how many iterations (matvecs) required?

- ▶ conjugate gradient
 - ▶ $O(\sqrt{\kappa} \log(\frac{1}{\epsilon}))$
- ▶ gradient descent (GD)
 - ▶ $O(\kappa \log(1/\epsilon))$
- ▶ accelerated gradient descent
 - ▶ $O\left(\sqrt{\kappa} \log\left(\frac{R^2}{\epsilon}\right)\right)$ more generalizable, but more parameters to tune

source: [KV16, Bub14]

Residual

define **residual** $r = b - Ax$ at putative solution x

► $r = -\nabla f(x) = A(x_\star) - x$

Residual

define **residual** $r = b - Ax$ at putative solution x

► $r = -\nabla f(x) = A(x_*) - x$

measures of error:

- objective function $f(x) - f(x_*)$
- norm of residual $\|r\|$
- norm of gradient $\|\nabla f(x)\|$
- in terms of r , can compute error in objective

$$f(x) - f(x_*) = \|x - x_*\|_A \quad (1)$$

$$= \frac{1}{2}(x - x_*)^T A(x - x_*) \quad (2)$$

$$= \frac{1}{2}(r)^T A^{-1}(r) \quad (3)$$

$$= \|r\|_{A^{-1}} \quad (4)$$

Krylov subspace

the Krylov subspace of dimension k is

$$\mathcal{K}_k = \text{span}\{b, Ab, \dots, A^{k-1}b\} = \text{span}\{p_k(A)b \mid \text{degree}(p) < k\}$$

Krylov subspace

the Krylov subspace of dimension k is

$$\mathcal{K}_k = \text{span}\{b, Ab, \dots, A^{k-1}b\} = \text{span}\{p_k(A)b \mid \text{degree}(p) < k\}$$

the iterates of the **Krylov sequence** $x^{(1)}, x^{(2)}, \dots$, minimize objective over successive Krylov subspaces

$$x^{(k)} = \underset{x \in \mathcal{K}_k}{\operatorname{argmin}} f(x) = \underset{x \in \mathcal{K}_k}{\operatorname{argmin}} \|Ax - b\| = \underset{x \in \mathcal{K}_k}{\operatorname{argmin}} \|x - x_\star\|_A$$

the CG algorithm generates the Krylov sequence

Properties of Krylov sequence

- ▶ $f(x^{(k+1)}) \leq f(x^{(k)})$ (but $\|r\|$ can increase)
- ▶ $x^{(n)} = x_*$
- ▶ $x^{(k)} = p_k(A)b$, where p_k is a polynomial with degree $< k$
- ▶ less obvious: there is a two-term recurrence

$$x^{(k+1)} = x^{(k)} + \alpha_k r^{(k)} + \beta_k (x^{(k)} - x^{(k-1)})$$

- ▶ α_k and β_k are determined by the CG algorithm
- ▶ looks like accelerated gradient descent update
- ▶ can derive recurrence from optimality conditions:
each new iterate $x^{(k+1)}$ must have gradient (residual)
orthogonal to \mathcal{K}_k

Properties of Krylov sequence

- ▶ $f(x^{(k+1)}) \leq f(x^{(k)})$ (but $\|r\|$ can increase)
- ▶ $x^{(n)} = x_*$
- ▶ $x^{(k)} = p_k(A)b$, where p_k is a polynomial with degree $< k$
- ▶ less obvious: there is a two-term recurrence

$$x^{(k+1)} = x^{(k)} + \alpha_k r^{(k)} + \beta_k (x^{(k)} - x^{(k-1)})$$

- ▶ α_k and β_k are determined by the CG algorithm
- ▶ looks like accelerated gradient descent update
- ▶ can derive recurrence from optimality conditions:
each new iterate $x^{(k+1)}$ must have gradient (residual)
orthogonal to \mathcal{K}_k

exercise: derive CG update from Krylov optimality condition

CG converges in $\text{Rank}(A)$ iterations

write (don't compute!) SVD of $A = V\Lambda V^T$ with

- ▶ $r = \text{Rank}(A)$
- ▶ $\Lambda \in \mathbf{R}^r \times r$ diagonal
- ▶ $V \in \mathbf{R}^{n \times r}$: orthonormal: $V^T V = I_r$

CG converges in $\text{Rank}(A)$ iterations

write (don't compute!) SVD of $A = V\Lambda V^T$ with

- ▶ $r = \text{Rank}(A)$
- ▶ $\Lambda \in \mathbf{R}^r \times r$ diagonal
- ▶ $V \in \mathbf{R}^{n \times r}$: orthonormal: $V^T V = I_r$

characteristic polynomial of Λ :

$$\xi(s) = \det(sI_r - \Lambda) = (s - \lambda_1) \cdots (s - \lambda_r) = s^r + \alpha s^{r-1} + \cdots + \alpha_r$$

Cayley-Hamilton theorem

$$\begin{aligned}\xi(\Lambda) = 0 &= \Lambda^r + \alpha_1 \Lambda^{r-1} + \cdots + \alpha_r I_r \\ \Lambda^{-1} &= -(1/\alpha_r)(\Lambda^{r-1} + \alpha_1 \Lambda^{r-2} + \cdots + \alpha_{r-1} I_r)\end{aligned}$$

CG converges in Rank(A) iterations

write (don't compute!) SVD of $A = V\Lambda V^T$ with

- ▶ $r = \text{Rank}(A)$
- ▶ $\Lambda \in \mathbf{R}^r \times r$ diagonal
- ▶ $V \in \mathbf{R}^{n \times r}$: orthonormal: $V^T V = I_r$

characteristic polynomial of Λ :

$$\xi(s) = \det(sI_r - \Lambda) = (s - \lambda_1) \cdots (s - \lambda_r) = s^r + \alpha s^{r-1} + \cdots + \alpha_r$$

Cayley-Hamilton theorem

$$\begin{aligned}\xi(\Lambda) = 0 &= \Lambda^r + \alpha_1 \Lambda^{r-1} + \cdots + \alpha_r I_r \\ \Lambda^{-1} &= -(1/\alpha_r)(\Lambda^{r-1} + \alpha_1 \Lambda^{r-2} + \cdots + \alpha_{r-1} I_r)\end{aligned}$$

write $A^{-1} = V\Lambda^{-1}V^T$ in terms of this decomposition:

$$\begin{aligned}A^{-1} = V\Lambda^{-1}V^T &= -(1/\alpha_r)(V\Lambda^{r-1}V^T + \alpha_1 V\Lambda^{r-2}V^T + \cdots + \alpha_{r-1} I) \\ &= -(1/\alpha_r)(A^{r-1} + \alpha_1 A^{r-2} + \cdots + \alpha_{r-1} I)\end{aligned}$$

in particular, $x_* = A^{-1}b \in \mathcal{K}_r$

Outline

QR

Conjugate gradient

Iterative refinement

Iterative refinement

want to solve $Ax = b$.

given approximate solution $Ax^{(0)} \approx b$, for $k = 1, \dots$,

- ▶ compute residual $r^{(k)} = b - Ax^{(k)}$
- ▶ use any method to solve $A\delta^{(k)} = r^{(k)}$
- ▶ $x^{(k+1)} = x^{(k)} + \delta^{(k)}$



Sébastien Bubeck.

Convex optimization: Algorithms and complexity, 2014.



Sahar Karimi and Stephen A. Vavasis.

A unified convergence bound for conjugate gradient and accelerated gradient, 2016.